

```
In [1]: import tensorflow as tf
from tensorflow.keras import models, layers
import matplotlib.pyplot as plt
```

```
In [2]: ## Loading the data directory --
dataset = tf.keras.preprocessing.image_dataset_from_directory(
    "..\\Data\\animals\\",
    batch_size=32,
    image_size=(224, 224),
    shuffle=True,
    color_mode="rgb"
)
```

Found 5400 files belonging to 90 classes.

```
In [3]: ## class_names
class_names = dataset.class_names
class_names
```

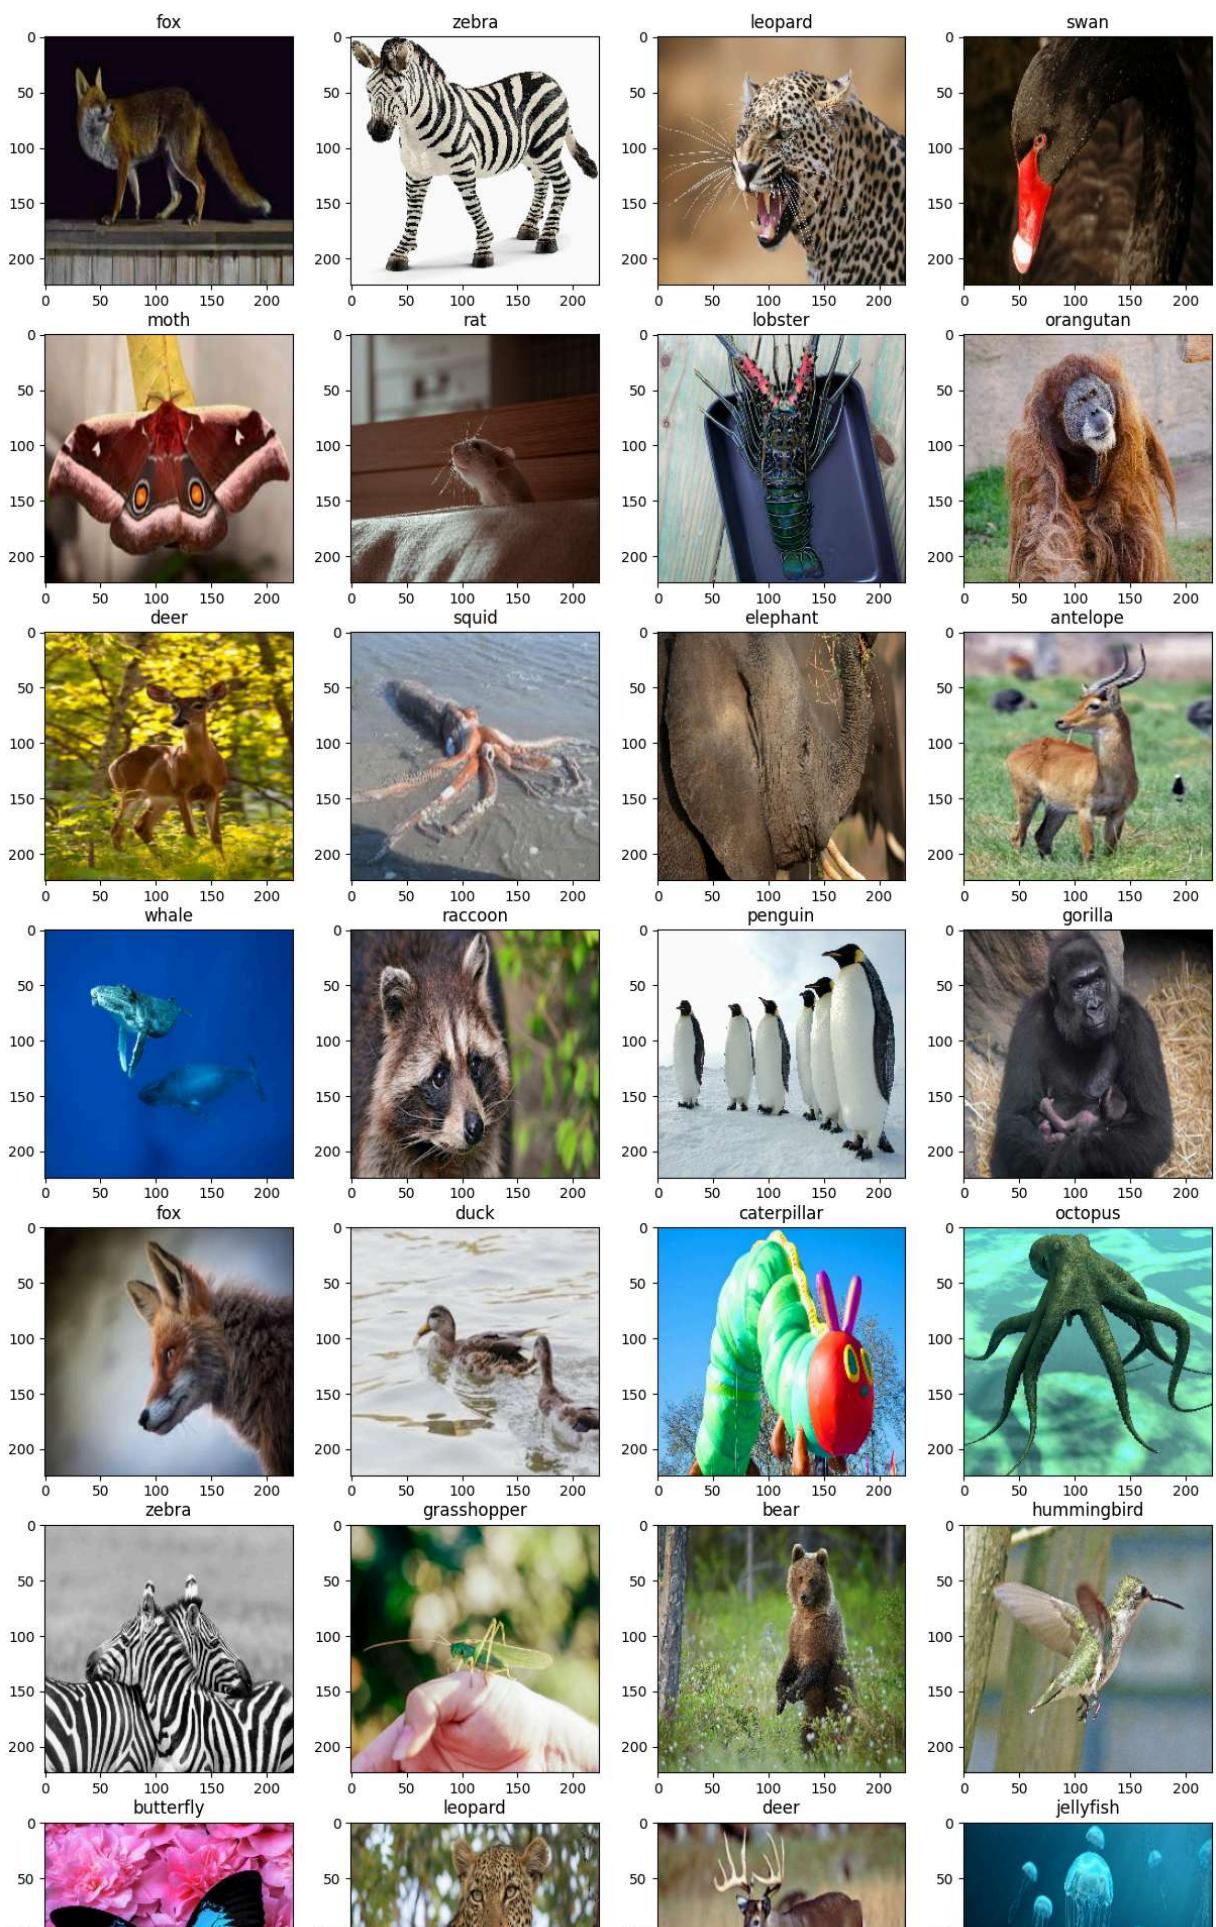
```
Out[3]: ['antelope',
 'badger',
 'bat',
 'bear',
 'bee',
 'beetle',
 'bison',
 'boar',
 'butterfly',
 'cat',
 'caterpillar',
 'chimpanzee',
 'cockroach',
 'cow',
 'coyote',
 'crab',
 'crow',
 'deer',
 'dog',
 'dolphin',
 'donkey',
 'dragonfly',
 'duck',
 'eagle',
 'elephant',
 'flamingo',
 'fly',
 'fox',
 'goat',
 'goldfish',
 'goose',
 'gorilla',
 'grasshopper',
 'hamster',
 'hare',
 'hedgehog',
 'hippopotamus',
 'hornbill',
 'horse',
 'hummingbird',
 'hyena',
 'jellyfish',
 'kangaroo',
 'koala',
 'ladybugs',
 'leopard',
 'lion',
 'lizard',
 'lobster',
 'mosquito',
 'moth',
 'mouse',
 'octopus',
 'okapi',
 'orangutan',
 'otter',
```

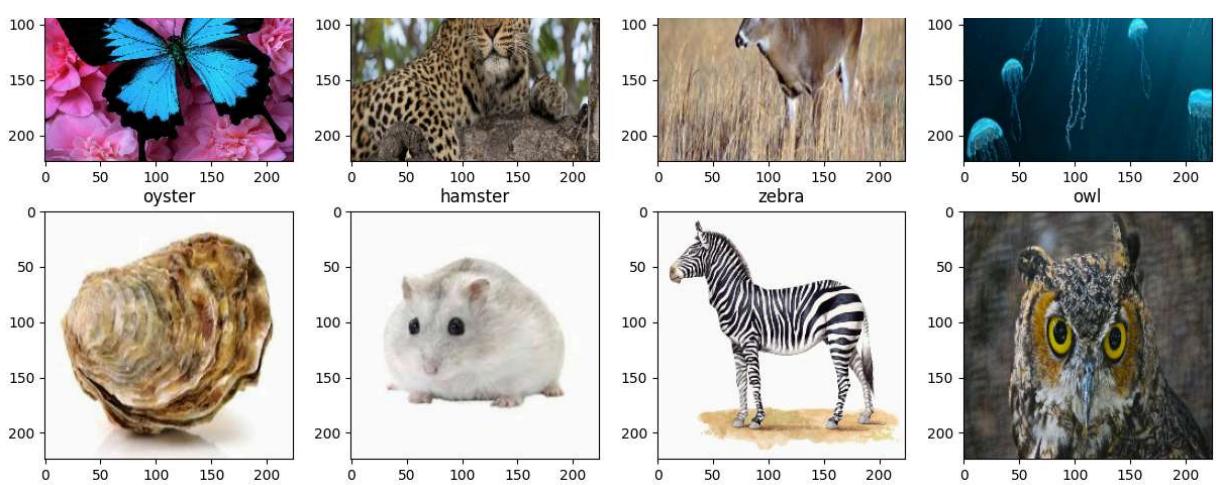
```
'owl',
'ox',
'oyster',
'panda',
'parrot',
'pelecaniformes',
'penguin',
'pig',
'pigeon',
'porcupine',
'possum',
'raccoon',
'rat',
'reindeer',
'rhinoceros',
'sandpiper',
'seahorse',
'seal',
'shark',
'sheep',
'snake',
'sparrow',
'squid',
'squirrel',
'starfish',
'swan',
'tiger',
'turkey',
'turtle',
'whale',
'wolf',
'wombat',
'woodpecker',
'zebra']
```

```
In [4]: len(dataset) ## 32 images per batch ==> 169*32=5408
```

```
Out[4]: 169
```

```
In [5]: plt.figure(figsize=(15,30))
for image_batch,label_batch in dataset.take(1):
    for i in range(32):
        plt.subplot(8,4,i+1)
        plt.imshow(image_batch[i].numpy().astype("uint8"))
        plt.title(class_names[label_batch[i]])
```





Splitting the Dataset into Train(80%), Test(10%) and Validation(10%)

```
In [6]: def get_dataset_partitions(ds,train_split=0.8,val_split=0.1,test_split=0.1,shuffle=True):
    ds_size = len(dataset)

    if shuffle:
        ds = ds.shuffle(shuffle_size,seed=12)
    train_size = int(train_split*ds_size)
    val_size = int(val_split*ds_size)

    train_ds = ds.take(train_size)
    validation_ds = ds.skip(train_size).take(val_size)
    test_ds = ds.skip(train_size).skip(val_size)

    return train_ds,validation_ds,test_ds
```

```
In [7]: train_ds, val_ds, test_ds = get_dataset_partitions(dataset)
```

```
In [8]: len(train_ds),len(test_ds),len(val_ds)
```

```
Out[8]: (135, 18, 16)
```

Resizing and Rescaling

```
In [9]: from tensorflow.keras.applications.vgg16 import preprocess_input
```

```
In [10]: ## Applying VGG16 preprocessing
```

```
def process(image,label):
    image = tf.cast(image,tf.float32)
    image = preprocess_input(image)
    return image,label
```

```
In [11]: ## Applying process function to the dataset
train_ds = train_ds.map(process).cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)
val_ds = val_ds.map(process).cache().prefetch(buffer_size=tf.data.AUTOTUNE)
test_ds = test_ds.map(process).cache().prefetch(buffer_size=tf.data.AUTOTUNE)

In [12]: resize_and_rescale = tf.keras.Sequential([
    layers.Resizing(224,224),
])

In [13]: data_augmentation = tf.keras.Sequential([
    layers.RandomFlip("horizontal_and_vertical"),
    layers.RandomRotation(0.2),
    layers.RandomZoom(height_factor=(0.2, 0.3)),
    layers.RandomZoom(height_factor=(-0.3,-0.2)),
])
```

Transfer Learning (Pretrained model)

```
In [14]: ## Loading VGG16 model

conv_base = tf.keras.applications.VGG16(
    include_top=False,
    weights="imagenet",
    input_shape=(224,224,3)
)
```

Fine Tuning

```
In [15]: ## Freezing the Convolution Layers and Training the Dense Layers and Some Convolutional Layers
conv_base.trainable = True

set_trainable = False

for layer in conv_base.layers:
    if layer.name == 'block5_conv1':
        set_trainable = True
    if set_trainable:
        layer.trainable = True
    else: layer.trainable = False
for layer in conv_base.layers:
    print(layer.name,layer.trainable)
```

```
input_layer False
block1_conv1 False
block1_conv2 False
block1_pool False
block2_conv1 False
block2_conv2 False
block2_pool False
block3_conv1 False
block3_conv2 False
block3_conv3 False
block3_pool False
block4_conv1 False
block4_conv2 False
block4_conv3 False
block4_pool False
block5_conv1 True
block5_conv2 True
block5_conv3 True
block5_pool True
```

```
In [16]: model = models.Sequential([
    resize_and_rescale,
    data_augmentation,
    conv_base,
    layers.GlobalAveragePooling2D(),
    layers.Dense(256, activation='relu'), ## Relu activation function
    layers.Dropout(0.3), ## Using Dropout to avoid overfitting
    layers.Dense(len(class_names), activation='softmax') ## As the dataset is a mult
])
```

```
In [17]: model.build(input_shape=(None, 224, 224, 3)) #force building the model
```

```
In [18]: model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	
sequential (Sequential)	(None, 224, 224, 3)	
sequential_1 (Sequential)	(None, 224, 224, 3)	
vgg16 (Functional)	(None, 7, 7, 512)	1
global_average_pooling2d (GlobalAveragePooling2D)	(None, 512)	
dense (Dense)	(None, 256)	
dropout (Dropout)	(None, 256)	
dense_1 (Dense)	(None, 90)	

```
Total params: 14,869,146 (56.72 MB)
Trainable params: 7,233,882 (27.60 MB)
Non-trainable params: 7,635,264 (29.13 MB)
```

```
In [19]: model.compile(
    optimizer=tf.optimizers.RMSprop(learning_rate=1e-4),
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy'])
## compiling the model at 0.0001 Learning rate
```

```
In [20]: from tensorflow.keras.callbacks import EarlyStopping
```

```
In [21]: early_stopping_callbacks = EarlyStopping(monitor='val_loss', patience=10, restore_bes
```

```
In [22]: history = model.fit(train_ds, validation_data=val_ds, epochs=25, callbacks=[early_stop
```

Epoch 1/25
135/135 408s 3s/step - accuracy: 0.0232 - loss: 4.5402 - val_accuracy: 0.0273 - val_loss: 4.4064

Epoch 2/25
135/135 388s 3s/step - accuracy: 0.0427 - loss: 4.3200 - val_accuracy: 0.0781 - val_loss: 4.0476

Epoch 3/25
135/135 384s 3s/step - accuracy: 0.1078 - loss: 4.0219 - val_accuracy: 0.1816 - val_loss: 3.5956

Epoch 4/25
135/135 386s 3s/step - accuracy: 0.1695 - loss: 3.6665 - val_accuracy: 0.2520 - val_loss: 3.2220

Epoch 5/25
135/135 385s 3s/step - accuracy: 0.2303 - loss: 3.3142 - val_accuracy: 0.3242 - val_loss: 2.7943

Epoch 6/25
135/135 378s 3s/step - accuracy: 0.3008 - loss: 2.9517 - val_accuracy: 0.4121 - val_loss: 2.3398

Epoch 7/25
135/135 395s 3s/step - accuracy: 0.3525 - loss: 2.6145 - val_accuracy: 0.5391 - val_loss: 1.9623

Epoch 8/25
135/135 364s 3s/step - accuracy: 0.4265 - loss: 2.2971 - val_accuracy: 0.5723 - val_loss: 1.7092

Epoch 9/25
135/135 359s 3s/step - accuracy: 0.4879 - loss: 2.0093 - val_accuracy: 0.6738 - val_loss: 1.3312

Epoch 10/25
135/135 357s 3s/step - accuracy: 0.5397 - loss: 1.8096 - val_accuracy: 0.6738 - val_loss: 1.3774

Epoch 11/25
135/135 357s 3s/step - accuracy: 0.5795 - loss: 1.6112 - val_accuracy: 0.7070 - val_loss: 1.0894

Epoch 12/25
135/135 357s 3s/step - accuracy: 0.6125 - loss: 1.4664 - val_accuracy: 0.7598 - val_loss: 0.9862

Epoch 13/25
135/135 357s 3s/step - accuracy: 0.6473 - loss: 1.3269 - val_accuracy: 0.7305 - val_loss: 1.0365

Epoch 14/25
135/135 355s 3s/step - accuracy: 0.6737 - loss: 1.1944 - val_accuracy: 0.7422 - val_loss: 1.0222

Epoch 15/25
135/135 356s 3s/step - accuracy: 0.7027 - loss: 1.0919 - val_accuracy: 0.7910 - val_loss: 0.8672

Epoch 16/25
135/135 356s 3s/step - accuracy: 0.7210 - loss: 1.0268 - val_accuracy: 0.8105 - val_loss: 0.8362

Epoch 17/25
135/135 356s 3s/step - accuracy: 0.7393 - loss: 0.9682 - val_accuracy: 0.8105 - val_loss: 0.9132

Epoch 18/25
135/135 357s 3s/step - accuracy: 0.7491 - loss: 0.9067 - val_accuracy: 0.8340 - val_loss: 0.7585

Epoch 19/25
135/135 353s 3s/step - accuracy: 0.7658 - loss: 0.8738 - val_accuracy: 0.8500 - val_loss: 0.7200

```

curacy: 0.8281 - val_loss: 0.6951
Epoch 20/25
135/135 356s 3s/step - accuracy: 0.7806 - loss: 0.7809 - val_ac
curacy: 0.8086 - val_loss: 0.8104
Epoch 21/25
135/135 354s 3s/step - accuracy: 0.7910 - loss: 0.7476 - val_ac
curacy: 0.8203 - val_loss: 0.9920
Epoch 22/25
135/135 356s 3s/step - accuracy: 0.8087 - loss: 0.7026 - val_ac
curacy: 0.8262 - val_loss: 0.9028
Epoch 23/25
135/135 381s 3s/step - accuracy: 0.8170 - loss: 0.6678 - val_ac
curacy: 0.8027 - val_loss: 0.8082
Epoch 24/25
135/135 365s 3s/step - accuracy: 0.8237 - loss: 0.6542 - val_ac
curacy: 0.8594 - val_loss: 0.7876
Epoch 25/25
135/135 362s 3s/step - accuracy: 0.8356 - loss: 0.5752 - val_ac
curacy: 0.8496 - val_loss: 0.8100

```

In [23]: `scores = model.evaluate(test_ds) ## Evaluating the model on test_data`

```
18/18 46s 2s/step - accuracy: 0.8559 - loss: 0.6677
```

Predictions on Test Data

In [31]:

```

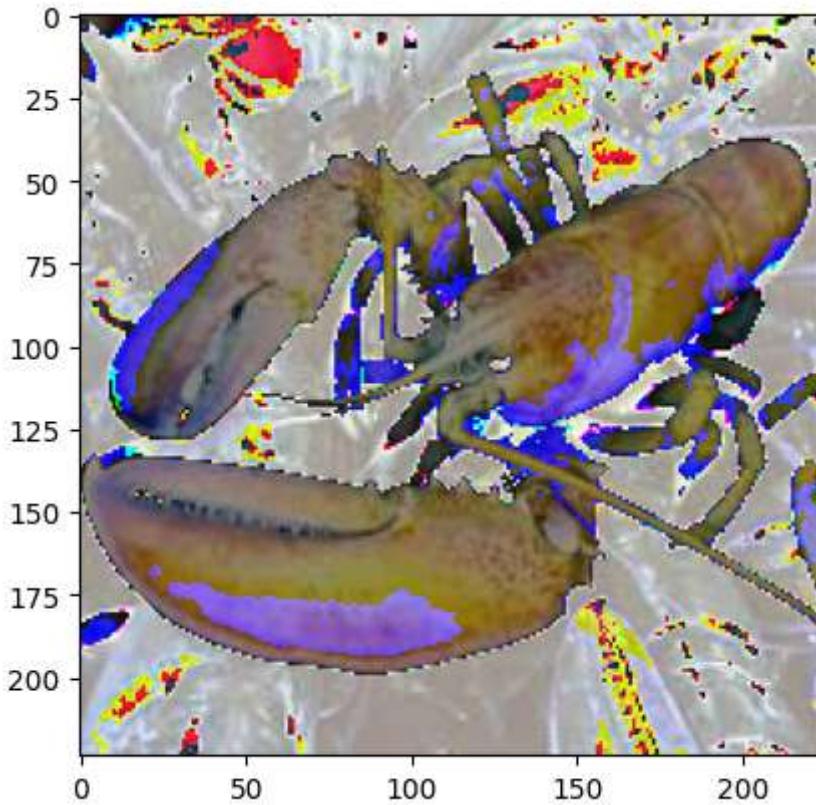
for image_batch, labels_batch in test_ds.take(1):
    first_image = image_batch[0].numpy().astype('uint8')
    first_label = labels_batch[0].numpy()

    print("first image prediction : ")
    plt.imshow(first_image)
    print("actual label:", class_names[first_label])

    batch_prediction = model.predict(image_batch)
    print("predicted label:", class_names[np.argmax(batch_prediction[0])])

```

first image prediction :
actual label: lobster
1/1 **3s** 3s/step
predicted label: lobster



```
In [26]: def predict(model,img):
    img_array = tf.keras.preprocessing.image.img_to_array(images[i].numpy())
    img_array = tf.expand_dims(img_array,0)

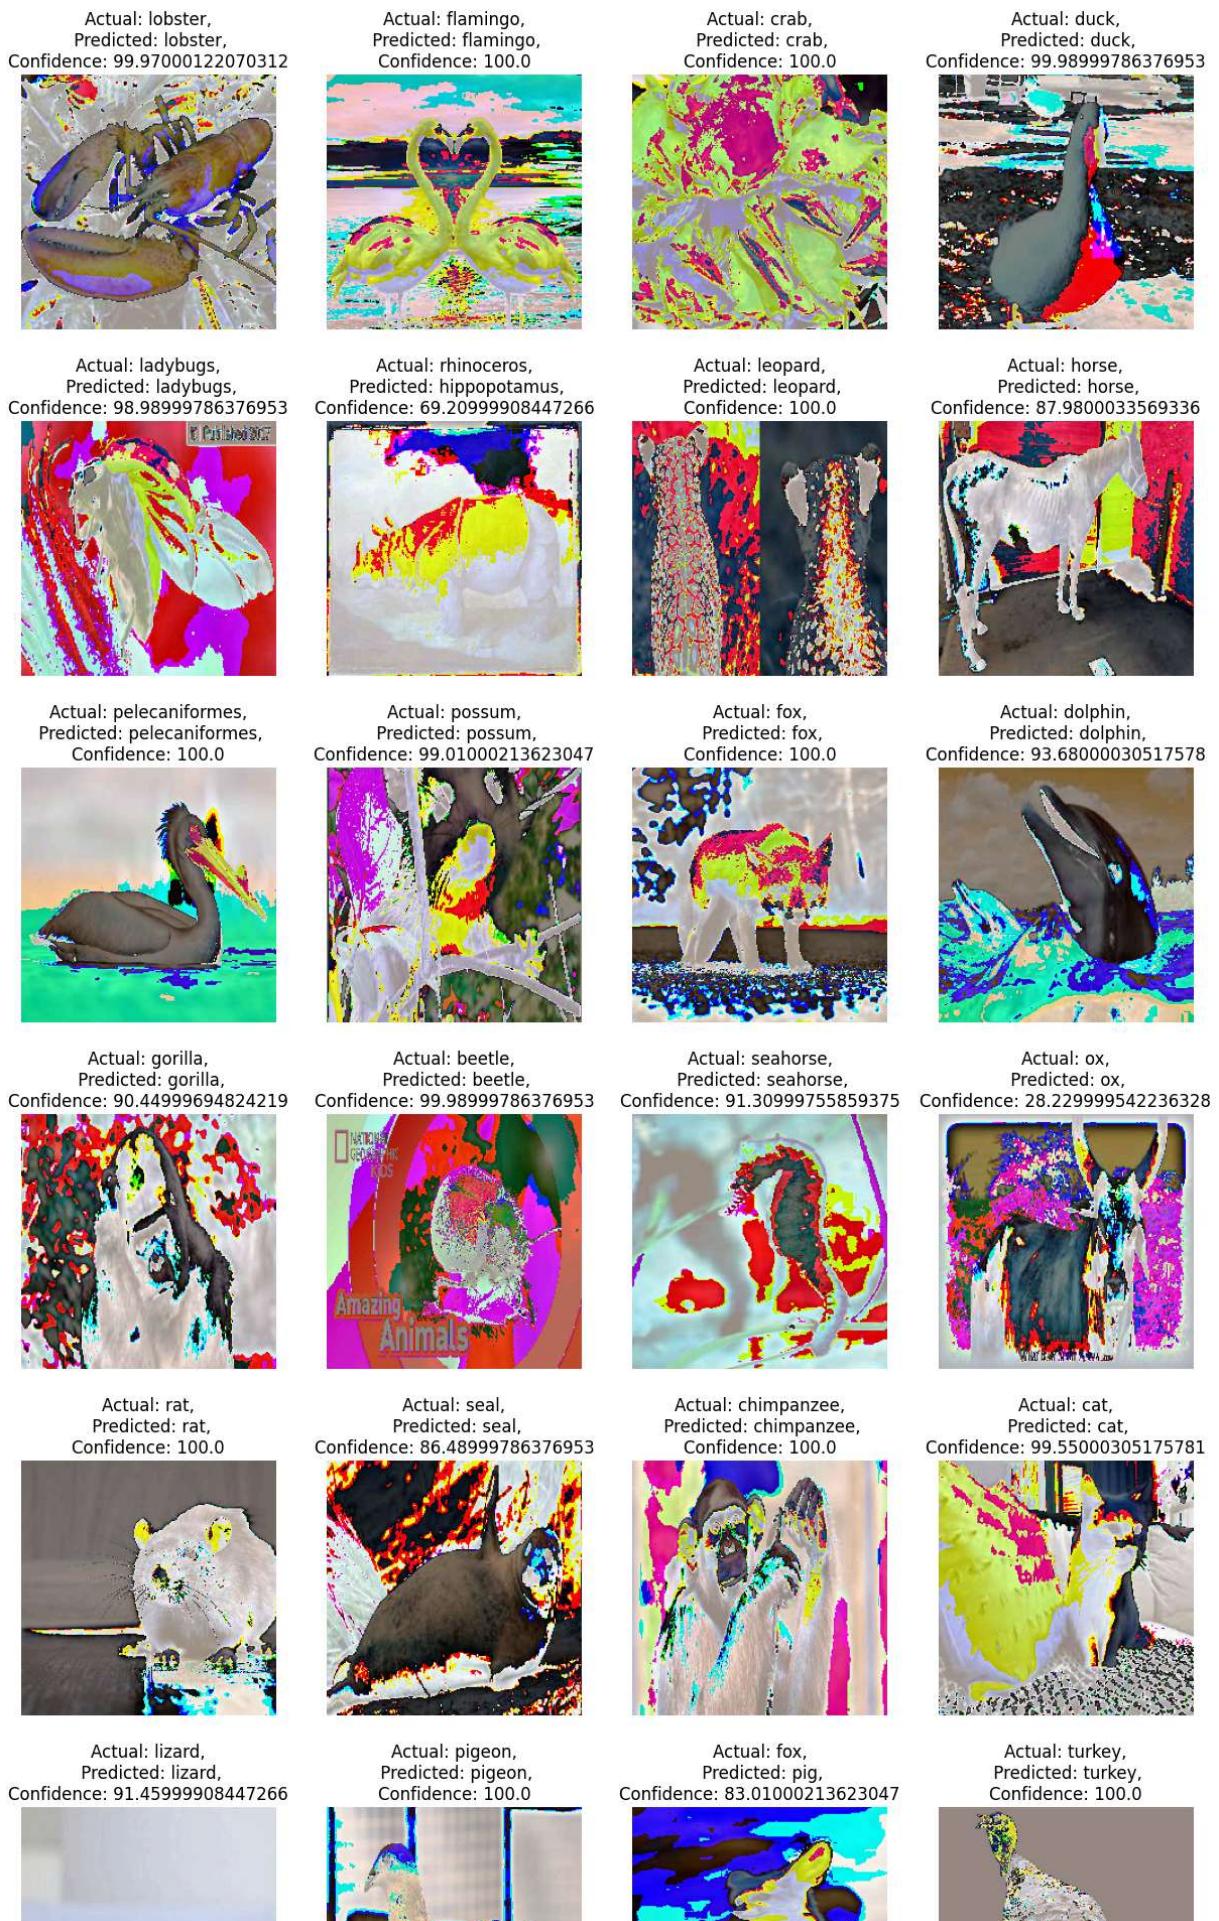
    predictions = model.predict(img_array)

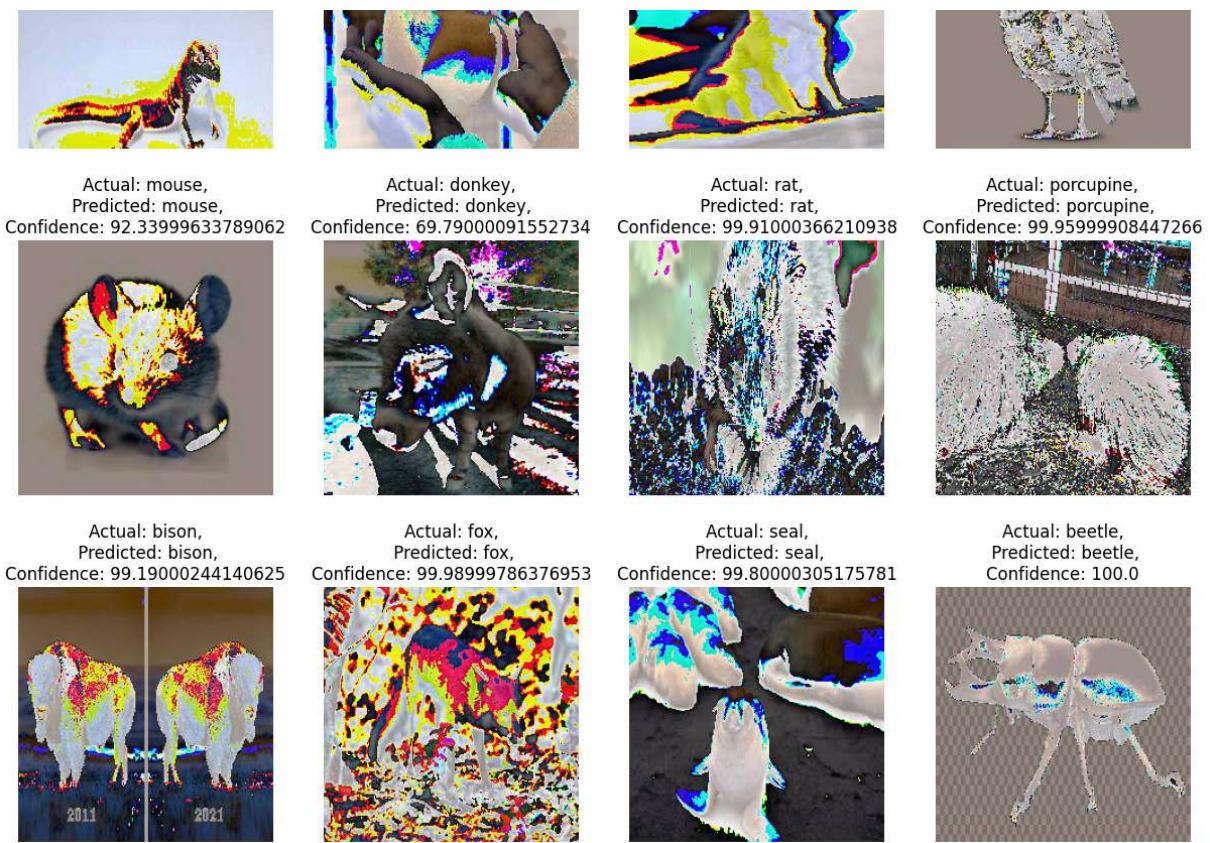
    predicted_class = class_names[np.argmax(predictions[0])]
    confidence = round(100*(np.max(predictions[0])),2)
    return predicted_class,confidence
```

```
In [27]: plt.figure(figsize=(15,35))
for images,labels in test_ds.take(1):
    for i in range(32):
        plt.subplot(8,4,i+1)
        plt.imshow(images[i].numpy().astype("uint8"))
        predicted_class,confidence = predict(model,images[i].numpy())
        actual_class = class_names[labels[i]]

        plt.title(f"Actual: {actual_class},\n Predicted: {predicted_class},\nConfid
```

1/1 ————— 1s 712ms/step
1/1 ————— 0s 211ms/step
1/1 ————— 0s 282ms/step
1/1 ————— 0s 330ms/step
1/1 ————— 0s 276ms/step
1/1 ————— 0s 223ms/step
1/1 ————— 0s 236ms/step
1/1 ————— 0s 213ms/step
1/1 ————— 0s 258ms/step
1/1 ————— 0s 197ms/step
1/1 ————— 0s 291ms/step
1/1 ————— 0s 232ms/step
1/1 ————— 0s 196ms/step
1/1 ————— 0s 211ms/step
1/1 ————— 0s 226ms/step
1/1 ————— 0s 239ms/step
1/1 ————— 0s 216ms/step
1/1 ————— 0s 191ms/step
1/1 ————— 0s 172ms/step
1/1 ————— 0s 215ms/step
1/1 ————— 0s 166ms/step
1/1 ————— 0s 165ms/step
1/1 ————— 0s 280ms/step
1/1 ————— 0s 187ms/step
1/1 ————— 0s 242ms/step
1/1 ————— 0s 169ms/step
1/1 ————— 0s 246ms/step
1/1 ————— 0s 157ms/step
1/1 ————— 0s 218ms/step
1/1 ————— 0s 156ms/step
1/1 ————— 0s 226ms/step
1/1 ————— 0s 161ms/step





Predicting on single image

```
In [28]: from tensorflow.keras.preprocessing import image
import numpy as np

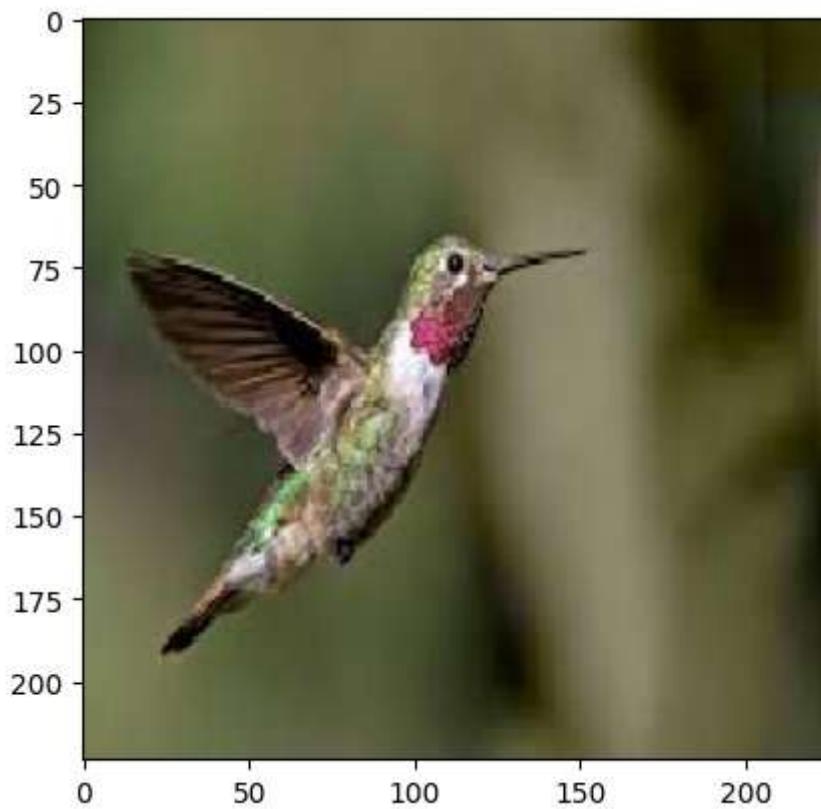
img_path = "D:/Artificial Intelligence/Deep Learning/Project2/Animal_Classifier/Dat

# Load and preprocess image
img = image.load_img(img_path, target_size=(224,224))
x = image.img_to_array(img)
plt.imshow(x.astype("uint8"))
x = np.expand_dims(x, axis=0)

# Applying VGG16 preprocessing
from tensorflow.keras.applications.vgg16 import preprocess_input
x = preprocess_input(x)

# Prediction
preds = model.predict(x)
pred_class = np.argmax(preds[0])
print("Predicted class:", class_names[pred_class])
```

1/1 ————— 0s 231ms/step
Predicted class: hummingbird



```
In [29]: # ## saving the model  
# model_version = 2  
# model.save(f"../models/classifier_{model_version}_acc_85.h5")
```

```
In [ ]:
```